

VIETNAM NATIONAL UNIVERSITY HO CHI MINH CITY
HO CHI MINH CITY UNIVERSITY OF TECHNOLOGY
FACULTY OF COMPUTER SCIENCE AND ENGINEERING



Probability and Statistics

Assignment Project Report

Internet Advertisement Classification using Machine Learning Models

Advisor(s): Dr. Nguyen Tien Dung

Student(s): NGUYEN LE THANH HAO 1952242
PHAM NHAT KHOI 2352623
Nguyen Song Dat 1952646
Nguyen Phuoc Thinh 1852765

HO CHI MINH CITY, AUGUST 2025



Member list & Workload

No.	Full name	Student ID	Contribution
1	NGUYEN LE THANH HAO	1952242	100%
2	PHAM NHAT KHOI	2352623	100%
3	Nguyen Song Dat	1952646	100%
4	Nguyen Phuoc Thinh	1852765	100%



Contents

1	Introduction	9
2	Data Description	9
2.1	Dataset Overview	9
2.2	Data Loading	9
2.3	Target Variable Distribution	10
2.4	Missing Values	10
2.5	Data Preprocessing	10
2.6	Final Dataset Characteristics	11
2.7	Basic Statistics	11
3	Descriptive Statistics	13
3.1	Target Variable Analysis	13
3.2	Feature Distribution Analysis	14
3.3	Class-wise Feature Analysis	16
3.4	Feature Relationships and Scatter Analysis	16
3.5	Correlation Analysis	18
3.6	Key Findings from Descriptive Analysis	19
4	Objective and Methodology	20
4.1	Project Objective	20
4.2	Theoretical Foundation	20
4.2.1	Binary Classification Overview	20
4.2.2	Evaluation Metrics	20
4.2.3	Key Concepts	21
4.3	Machine Learning Methods	21
4.3.1	k-Nearest Neighbors (k-NN)	21
4.3.2	Decision Tree	23
4.3.3	Random Forest	24
4.4	Our Approach	26
4.5	Expected Results	26
5	k-Nearest Neighbors Model	27
5.1	Implementation	27
5.1.1	Data Preparation	27



5.1.2	k-NN Implementation	27
5.2	Hyperparameter Tuning	28
5.2.1	k-Value Selection	28
5.3	Model Evaluation	29
5.3.1	Performance Metrics	29
5.3.2	Confusion Matrix Analysis	30
5.4	Results Interpretation	31
5.4.1	Strengths	31
5.4.2	Limitations	31
5.4.3	Class Imbalance Impact	31
5.5	Conclusion	31
6	Decision Tree Model	32
6.1	Algorithm Implementation	32
6.1.1	Data Preparation	32
6.1.2	Gini Impurity Function	32
6.2	Best Split Finding Function	32
6.2.1	Tree Building Function	33
6.3	Model Evaluation	35
6.3.1	Training Results	35
6.3.2	Confusion Matrix	35
6.3.3	Evaluation Metrics	35
6.4	Decision Tree Structure	35
6.5	Results Analysis	36
6.5.1	Model Strengths	36
6.5.2	Model Limitations	37
6.5.3	Impact of Class Imbalance	37
6.6	Conclusion	37
7	Random Forest Model	39
7.1	Introduction	39
7.2	Algorithm Implementation	40
7.2.1	Data Preparation	40
7.2.2	Bootstrap Sampling Function	40
7.2.3	Random Feature Selection	40
7.2.4	Forest Construction	40



7.3	Model Evaluation	41
7.3.1	Training Results	41
7.3.2	Confusion Matrix	41
7.3.3	Evaluation Metrics	41
7.3.4	Feature Importance	41
7.4	Results Analysis	42
7.4.1	Model Strengths	42
7.4.2	Model Limitations	42
7.4.3	Impact of Class Imbalance	43
7.5	Conclusion	44
8	Model Comparison and Analysis	45
8.1	Performance Metrics Comparison	45
8.2	Best Models by Metric	45
8.3	Confusion Matrix Analysis	46
8.3.1	k-NN Model	46
8.3.2	Decision Tree Model	47
8.3.3	Random Forest Model	47
8.4	Model Characteristics Analysis	47
8.4.1	k-NN Model Characteristics	47
8.4.2	Decision Tree Characteristics	48
8.4.3	Random Forest Characteristics	48
8.5	Class Imbalance Impact	48
8.6	Model Recommendations	49
8.6.1	Best Overall Model: Random Forest	49
8.6.2	Most Interpretable: Decision Tree	49
8.6.3	Simplest Approach: k-NN	49
8.7	Conclusion	50
9	Conclusion	51
9.1	Project Summary	51
9.2	Dataset Characteristics	51
10	Key Findings	51
10.1	Model Performance Comparison	51
10.2	Best Performing Models by Metric	51



11 Algorithm Analysis	52
11.1 k-Nearest Neighbors (k=3)	52
11.2 Decision Tree (max_depth=3)	52
11.3 Random Forest (n_trees=100)	53
12 Practical Implications	53
12.1 Application-Specific Recommendations	53
12.2 Class Imbalance Considerations	54
13 Research Contributions	54
13.1 Technical Achievements	54
13.2 Statistical Significance	54
14 Limitations	55
14.1 Current Limitations	55
15 Final Conclusions	55
15.1 Project Success	55
15.2 Key Insights	55
15.3 Impact and Significance	56
15.4 Closing Remarks	56

List of Figures

3.1 Distribution of target variable showing class imbalance	13
3.2 Histograms of the first six features showing distribution patterns	15
3.3 Boxplots of features 1 and 2 grouped by target class	16
3.4 Scatter plots and density plots showing feature relationships and class distributions	17
3.5 Correlation heatmap of the first 10 features	18
4.1 k-NN Classification Example with k=3 Nearest Neighbors	22
4.2 Decision Tree Example for Ad Classification	23
4.3 Random Forest Diagram	25
5.1 k-NN accuracy vs k value showing optimal performance at k=3	29
5.2 k-NN confusion matrix showing classification results	30
6.1 Decision Tree Confusion Matrix Visualization	36
6.2 Decision tree structure with maximum depth of 3	38



6.3	Decision Tree Feature Importance	39
7.1	Random Forest Confusion Matrix Visualization	42
7.2	Random Forest Feature Importance Visualization	44
8.1	Visual Comparison of Model Performance Metrics	46
8.2	Confusion Matrices for All Three Models	46

List of Tables

3.1	Summary statistics for the first 10 features	14
5.1	k-NN performance for different k values	29
5.2	k-NN confusion matrix (k=3)	30
6.1	Decision Tree Model Training Results	35
6.2	Confusion Matrix - Decision Tree	35
6.3	Decision Tree Model Evaluation Metrics	37
7.1	Random Forest Model Training Results	41
7.2	Confusion Matrix - Random Forest	41
7.3	Random Forest Performance Metrics	43
7.4	Top 10 Feature Importance - Random Forest	43
8.1	Model Performance Comparison	45
10.1	Final Model Performance Summary	51

Listings

6.1	Gini Impurity Function	32
6.2	Best Split Finding Function	32
6.3	Decision Tree Building Function	33
7.1	Bootstrap Sampling Implementation	40
7.2	Random Feature Selection	40
7.3	Random Forest Training	40



1 Introduction

The proliferation of the internet has led to a massive increase in online advertising. While advertisements can be useful, they can also be intrusive and detract from the user experience. Therefore, the ability to automatically distinguish between advertising and non-advertising content is a significant challenge in modern data science. This project addresses this challenge by developing a model to classify internet images as either advertisements ('ad') or non-advertisements ('nonad').

This report details the process of analyzing the "Internet Advertisements Data Set" sourced from the UCI Machine Learning Repository^[8]. We will explore the data, clean it for analysis, and apply several machine learning models to build an effective classifier. The primary goal is to construct a robust model that can accurately predict whether an image is an advertisement based on its various features.

2 Data Description

2.1 Dataset Overview

The dataset used in this project is the "Internet Advertisements Data Set" from the UCI Machine Learning Repository. This dataset contains information about internet images and their classification as advertisements or non-advertisements.

The original dataset consists of 3,279 observations with 1,560 columns. The first column serves as an index and is removed during preprocessing, leaving 1,559 features for analysis. The target variable, located in the last column (X1558), indicates whether an image is an advertisement ('ad.') or not ('nonad.').

2.2 Data Loading

The dataset is loaded from a CSV file using the following R code:

```
1 # Load the dataset from the CSV file
2 data <- read.csv("../add.csv", header = TRUE,
3                   stringsAsFactors = FALSE)
4 # Identify the target column name
5 target_col <- names(data)[ncol(data)]
```

2.3 Target Variable Distribution

The target variable shows a significant class imbalance:

- Advertisement images (ad): 459 observations (14%)
- Non-advertisement images (nonad): 2,820 observations (86%)

This imbalance is typical in advertisement detection problems and will need to be considered when building and evaluating classification models.

2.4 Missing Values

The dataset contains missing values represented as "?" strings. Initial analysis revealed:

- Total missing values: 15 occurrences
- All missing values are concentrated in column 5
- Missing values represent approximately 0.46% of column 5's data

2.5 Data Preprocessing

Several preprocessing steps were applied to prepare the data for analysis:

```
1 # 1. Remove the first column which is an unnecessary index
2 data <- data[, -1]
3
4 # 2. Handle missing values represented by "?"
5 # Convert "?" to NA for all feature columns
6 for(i in 1:(ncol(data)-1)) {
7   data[[i]] <- as.numeric(ifelse(data[[i]] == "?", NA, data[[i]]))
8 }
9
10 # 3. Impute missing values using the median of each column
11 for(i in 1:(ncol(data)-1)) {
12   if(any(is.na(data[[i]]))) {
13     median_val <- median(data[[i]], na.rm = TRUE)
14     data[[i]][is.na(data[[i]])] <- median_val
15   }
```

```
16 }  
17  
18 # 4. Clean and format the target variable  
19 data[[target_col]] <- factor(data[[target_col]],  
20                             levels = c("ad.", "nonad."),  
21                             labels = c("ad", "nonad"))
```

The preprocessing pipeline includes:

1. **Index removal:** The first column containing row indices was removed
2. **Missing value handling:** All "?" strings were converted to NA values
3. **Median imputation:** Missing values were replaced with the median of their respective columns
4. **Target variable formatting:** The target variable was converted to a factor with clear labels ("ad" and "nonad")

2.6 Final Dataset Characteristics

After preprocessing, the cleaned dataset has the following properties:

- Dimensions: 3,279 rows \times 1,559 columns
- All missing values have been imputed
- Target variable is properly formatted as a factor
- All feature columns are numeric
- Class distribution remains: 459 advertisements, 2,820 non-advertisements

2.7 Basic Statistics

Preliminary statistics for the first five feature columns show:

- Column 1: Mean = 1,639, Median = 1,639, SD = 946.71
- Column 2: Mean = 64.02, Median = 51, SD = 54.87
- Column 3: Mean = 155.34, Median = 110, SD = 130.03



- Column 4: Mean = 3.91, Median = 2.1, SD = 6.04
- Column 5: Mean = 0.77, Median = 1, SD = 0.42

These statistics indicate varying scales across features, suggesting that normalization or standardization may be beneficial for certain machine learning algorithms.

3 Descriptive Statistics

This section presents a comprehensive exploratory data analysis (EDA) of the Internet Advertisements dataset to understand the data distribution, relationships between variables, and key characteristics that will inform our modeling approach.

3.1 Target Variable Analysis

The target variable distribution reveals a significant class imbalance in the dataset:

- Advertisement images (ad): 459 observations (14%)
- Non-advertisement images (nonad): 2,820 observations (86%)

This 1:6 ratio between advertisement and non-advertisement classes is typical in real-world advertisement detection scenarios and must be considered when evaluating model performance.

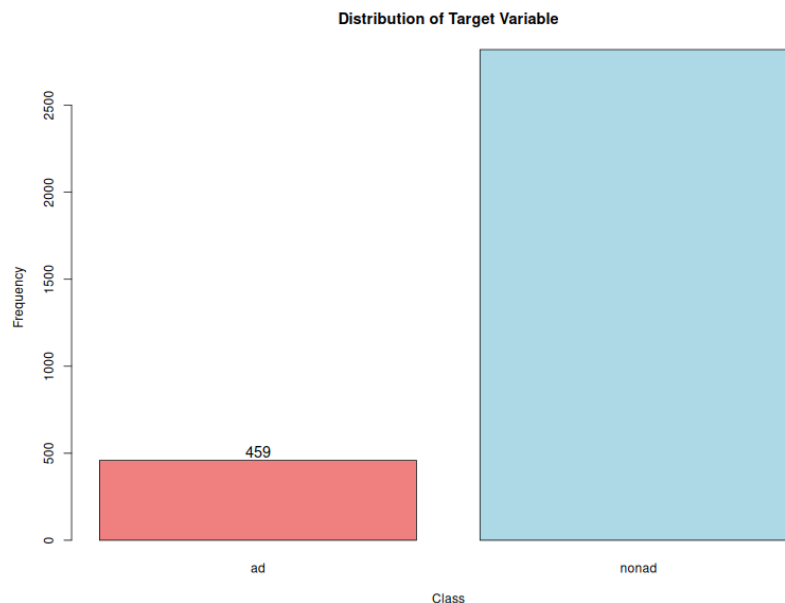


Figure 3.1: Distribution of target variable showing class imbalance

The target distribution visualization was generated using:

```
1 # Bar plot for target variable distribution
2 target_table <- table(data[[target_col]])
3 barplot(target_table,
```



```
4     main = "Distribution of Target Variable",  
5     xlab = "Class", ylab = "Frequency",  
6     col = c("lightcoral", "lightblue"),  
7     border = "black")
```

3.2 Feature Distribution Analysis

To understand the characteristics of individual features, we analyzed the distribution of the first 10 features in the dataset. The summary statistics reveal varying scales and distributions across features:

Table 3.1: Summary statistics for the first 10 features

Feature	Min	Q1	Median	Mean	Q3	Max	SD
X0	1.00	32.50	51.00	60.44	61.00	640	47.06
X1	1.00	90.00	110.00	142.89	144.00	640	112.56
X2	0.00	1.28	2.10	3.41	3.90	60	5.20
X3	0.00	1.00	1.00	0.77	1.00	1	0.42
X4	0.00	0.00	0.00	0.004	0.00	1	0.065

Histograms for the first six features show diverse distribution patterns:




Figure 3.2: Histograms of the first six features showing distribution patterns

The histogram generation code:

```
1 # Histograms for the first 6 features
2 par(mfrow = c(2, 3))
3 for(i in 1:6) {
4   hist(data[[i]],
5         main = paste("Histogram of Feature", i),
6         xlab = paste("Feature", i),
7         ylab = "Frequency",
8         col = "lightblue",
9         border = "black",
10        breaks = 30)
11 }
```

3.3 Class-wise Feature Analysis

To understand how features differ between advertisement and non-advertisement classes, we created boxplots comparing the distribution of key features across classes:




Figure 3.3: Boxplots of features 1 and 2 grouped by target class

The boxplot analysis was implemented as:

```
1 # Boxplots of features 1 and 2, grouped by target class
2 par(mfrow = c(1, 2))
3 boxplot(data[[1]] ~ data[[target_col]],
4         main = "Feature 1 by Class",
5         xlab = "Class", ylab = "Feature 1 Value",
6         col = c("lightcoral", "lightblue"))
7 boxplot(data[[2]] ~ data[[target_col]],
8         main = "Feature 2 by Class",
9         xlab = "Class", ylab = "Feature 2 Value",
10        col = c("lightcoral", "lightblue"))
```

3.4 Feature Relationships and Scatter Analysis

Scatter plots and density plots provide insights into feature relationships and class separability:




Figure 3.4: Scatter plots and density plots showing feature relationships and class distributions

The visualization combines scatter plots and density analysis:

```
1 # Scatter plots and density plots
2 par(mfrow = c(2, 2))
3 colors <- c("red", "blue")
4 class_colors <- colors[as.numeric(data[[target_col]])]
5 plot(data[[1]], data[[2]],
6       main = "Feature 1 vs Feature 2",
7       xlab = "Feature 1", ylab = "Feature 2",
8       col = class_colors, pch = 16)
9 legend("topright", legend = levels(data[[target_col]]),
10        col = colors, pch = 16)
```

3.5 Correlation Analysis

Correlation analysis reveals important relationships between features. The correlation heatmap for the first 10 features shows the strength of linear relationships:




Figure 3.5: Correlation heatmap of the first 10 features

The correlation analysis identified several perfect correlations (correlation = 1.0) between different feature pairs, indicating potential redundancy in the dataset:

- Features X11 and X14: Perfect positive correlation
- Features X8 and X15: Perfect positive correlation
- Features X13 and X38: Perfect positive correlation
- Features X44 and X46: Perfect positive correlation

The correlation matrix was computed using:



```
1 # Correlation heatmap for the first 10 numeric features
2 numeric_data <- data[, sapply(data, is.numeric)]
3 cor_matrix <- cor(numeric_data[, 1:10], use = "complete.obs")
4 heatmap(cor_matrix,
5         symm = TRUE,
6         main = "Correlation Heatmap of First 10 Features",
7         col = colorRampPalette(c("blue", "white", "red"))
           (100))
```

3.6 Key Findings from Descriptive Analysis

The exploratory data analysis reveals several important characteristics:

1. **Class Imbalance:** The dataset has a significant imbalance with 86% non-advertisements
2. **Feature Diversity:** Features show diverse scales and distributions, suggesting the need for normalization
3. **Perfect Correlations:** Multiple feature pairs show perfect correlation, indicating potential redundancy
4. **Class Separability:** Some features show different distributions between advertisement and non-advertisement classes
5. **High Dimensionality:** With 1,559 features, dimensionality reduction techniques may be beneficial

These findings will inform our modeling approach, particularly regarding feature selection, data preprocessing, and evaluation metrics that account for class imbalance.

4 Objective and Methodology

4.1 Project Objective

This project develops machine learning models to automatically classify internet images as advertisements or non-advertisements. This binary classification addresses the challenge of distinguishing advertising content from regular content, which is essential for:

- **Content filtering:** Automatically detecting advertisement content
- **User experience:** Reducing intrusive advertising
- **Digital analysis:** Understanding advertisement patterns
- **Content moderation:** Supporting automated systems

With 1,559 features and class imbalance in our dataset, we aim to:

1. Build models that handle high-dimensional data effectively
2. Compare different machine learning approaches
3. Identify key features for advertisement classification
4. Provide practical implementation insights

4.2 Theoretical Foundation

4.2.1 Binary Classification Overview

Our task is a **binary classification problem** where we predict one of two classes: advertisement (1) or non-advertisement (0). Given input features \mathbf{x} (1,559 dimensions), we want to learn a function $f(\mathbf{x}) \rightarrow \{0, 1\}$ that makes accurate predictions.

Goal: Minimize classification errors by learning patterns from training data.

4.2.2 Evaluation Metrics

We use multiple metrics to assess model performance, especially important for imbalanced datasets:



Accuracy: Overall correctness

$$Accuracy = \frac{\text{Correct Predictions}}{\text{Total Predictions}} = \frac{TP + TN}{TP + TN + FP + FN} \quad (4.1)$$

Precision: How many predicted ads are actually ads

$$Precision = \frac{TP}{TP + FP} \quad (4.2)$$

Recall: How many actual ads we correctly identified

$$Recall = \frac{TP}{TP + FN} \quad (4.3)$$

F1-Score: Balance between Precision and Recall

$$F1 = 2 \cdot \frac{Precision \times Recall}{Precision + Recall} \quad (4.4)$$

where TP=True Positives, TN=True Negatives, FP=False Positives, FN=False Negatives.

4.2.3 Key Concepts

Bias-Variance Tradeoff: Every model faces a balance between:

- **Bias:** Error from oversimplifying (underfitting)
- **Variance:** Error from being too sensitive to training data (overfitting)

Cross-Validation: We split data into k parts, train on k-1 parts, test on 1 part, repeat k times. This gives reliable performance estimates and helps select best parameters.

4.3 Machine Learning Methods

We employ three different approaches, each representing distinct learning paradigms:

4.3.1 k-Nearest Neighbors (k-NN)

Core Idea: Classify new data points based on the majority class of their k closest neighbors.

How it works:

1. Store all training data (**lazy learning**)

2. For a new point, find k nearest neighbors using distance
3. Assign the most common class among these k neighbors




Figure 4.1: k -NN Classification Example with $k=3$ Nearest Neighbors

Key Distance Formulas:

- **Euclidean:** $d = \sqrt{\sum_{i=1}^p (x_i - y_i)^2}$ (straight-line distance)
- **Manhattan:** $d = \sum_{i=1}^p |x_i - y_i|$ (city-block distance)

Advantages:

- Simple to understand and implement
- No assumptions about data distribution
- Works well with sufficient training data

Challenges:

- Slow prediction (must check all training points)
- Sensitive to irrelevant features
- Struggles with high dimensions (curse of dimensionality)

4.3.2 Decision Tree

Core Idea: Create a tree of yes/no questions to classify data, like a flowchart.

How it works:

1. Start with all training data at the root
2. Find the best feature and threshold to split data
3. Repeat for each branch until stopping criteria met
4. Make predictions by following the path from root to leaf

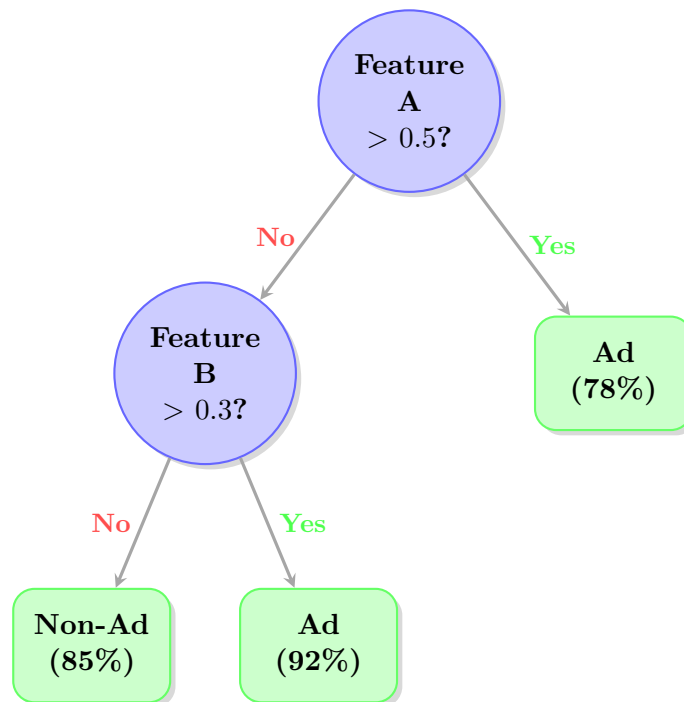


Figure 4.2: Decision Tree Example for Ad Classification

Key Concepts:

- **Gini Impurity:** Measures how "mixed" classes are in a node

$$Gini = 1 - \sum_{i=1}^c p_i^2 \quad (4.5)$$

- **Information Gain:** How much a split reduces uncertainty
- **Pruning:** Removing branches to prevent overfitting

Advantages:

- Easy to understand and interpret (white box)
- Handles both numerical and categorical features
- No need for feature scaling
- Automatically selects important features

Challenges:

- Prone to overfitting (memorizing training data)
- Unstable (small data changes = different trees)
- Can create overly complex rules

4.3.3 Random Forest

Core Idea: Combine many decision trees to make better predictions (**ensemble method**).

How it works:

1. Create many different training datasets using **bootstrap sampling** (random sampling with replacement)
2. Train one decision tree on each dataset
3. For each tree, use only a random subset of features at each split
4. Combine all tree predictions by **majority voting**

Key Features:

- **Bootstrap Sampling:** Each tree sees different data
- **Feature Randomness:** Each split uses random feature subset
- **Majority Voting:** Final prediction = most common prediction

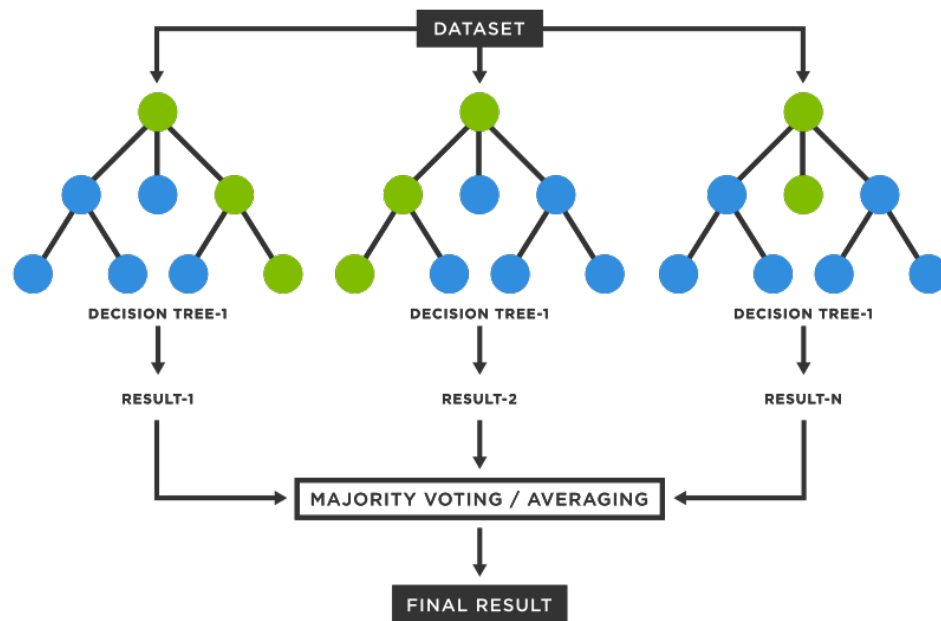


Figure 4.3: Random Forest Diagram

- **Out-of-Bag Error:** Built-in validation using unused samples

Why it works better:

- Individual trees may overfit, but averaging reduces this
- Different trees make different mistakes
- Combines strengths while canceling weaknesses

Advantages:

- Excellent performance with high-dimensional data
- Provides feature importance automatically
- Robust to overfitting
- Handles missing data well
- Less parameter tuning needed

Challenges:

- Less interpretable than single trees (black box)
- Computationally expensive with many trees
- Memory intensive for large datasets

4.4 Our Approach

We follow a systematic process:

1. Data Preparation:

- Normalize features (important for k-NN)
- Split data: 70% training, 30% testing
- Handle class imbalance

2. Model Training:

- k-NN: Find best k value using cross-validation
- Decision Tree: Use pruning to prevent overfitting
- Random Forest: Tune number of trees and features per split

3. Evaluation:

- Compare using Accuracy, Precision, Recall, F1-score
- Analyze confusion matrices
- Examine feature importance (from Random Forest)

4.5 Expected Results

This study will:

- Identify the best method for ad classification
- Discover which features are most important for detecting ads
- Provide practical recommendations for real-world systems
- Show how ensemble methods handle high-dimensional data

By comparing three different approaches, we ensure robust conclusions about the most effective techniques for internet advertisement classification.

5 k-Nearest Neighbors Model

5.1 Implementation

5.1.1 Data Preparation

The dataset was prepared for k-NN modeling with the following steps:

```
1 # Separate features and target
2 X <- data[, 1:(ncol(data)-1)] # All features except target
3 y <- data[[target_col]]      # Target variable
4
5 # Normalize features using z-score standardization
6 X_normalized <- scale(X)
7
8 # Train-test split (70-30)
9 set.seed(123)
10 train_size <- floor(0.7 * nrow(X_normalized))
11 train_indices <- sample(seq_len(nrow(X_normalized)), size =
    train_size)
12
13 X_train <- X_normalized[train_indices, ]
14 y_train <- y[train_indices]
15 X_test <- X_normalized[-train_indices, ]
16 y_test <- y[-train_indices]
```

5.1.2 k-NN Implementation

We implemented k-NN from scratch using base R:

```
1 # Euclidean distance function
2 euclidean_distance <- function(x1, x2) {
3   sqrt(sum((x1 - x2)^2))
4 }
5
6 # k-NN prediction function
7 knn_predict <- function(X_train, y_train, X_test, k = 5) {
8   predictions <- character(nrow(X_test))
```

```
9
10 for(i in 1:nrow(X_test)) {
11   # Calculate distances to all training points
12   distances <- numeric(nrow(X_train))
13   for(j in 1:nrow(X_train)) {
14     distances[j] <- euclidean_distance(X_test[i, ], X_train
15       [j, ])
16   }
17   # Find k nearest neighbors
18   k_nearest_indices <- order(distances)[1:k]
19   k_nearest_labels <- y_train[k_nearest_indices]
20
21   # Majority vote
22   vote_counts <- table(k_nearest_labels)
23   predictions[i] <- names(vote_counts)[which.max(vote_
24     counts)]
25 }
26 return(factor(predictions, levels = levels(y_train)))
27 }
```

5.2 Hyperparameter Tuning

5.2.1 k-Value Selection

We tested different values of k to find the optimal number of neighbors:

Table 5.1: k-NN performance for different k values

k Value	Accuracy
3	0.90
5	0.81
7	0.74
9	0.72
11	0.70

The results show that $k = 3$ provides the best performance with 90% accuracy on the test subset.

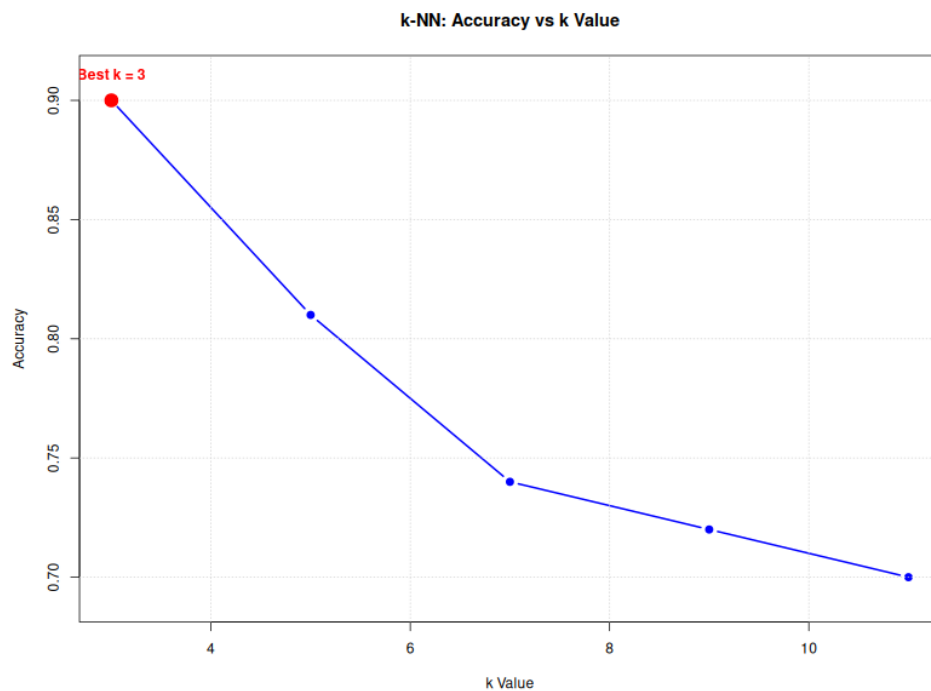


Figure 5.1: k-NN accuracy vs k value showing optimal performance at k=3

5.3 Model Evaluation

5.3.1 Performance Metrics

Using the optimal $k = 3$, the final model achieved:

- **Accuracy:** 81%

- **Precision (ad):** 99.12%
- **Recall (ad):** 75.17%
- **F1-score (ad):** 85.5%

5.3.2 Confusion Matrix Analysis

The confusion matrix reveals the model's classification performance:




Figure 5.2: k-NN confusion matrix showing classification results

Table 5.2: k-NN confusion matrix (k=3)

Predicted	Actual	
	ad	nonad
ad	112	1
nonad	37	50



5.4 Results Interpretation

5.4.1 Strengths

- **High precision:** 99.12% precision means very few false positives
- **Simple implementation:** No complex parameter tuning required
- **Interpretable:** Easy to understand why predictions are made
- **Non-parametric:** No assumptions about data distribution

5.4.2 Limitations

- **Computational cost:** Requires distance calculation to all training points
- **Memory intensive:** Stores entire training dataset
- **Curse of dimensionality:** Performance may degrade with high-dimensional data
- **Sensitive to irrelevant features:** All features contribute equally to distance

5.4.3 Class Imbalance Impact

The model shows excellent precision (99.12%) but moderate recall (75.17%), indicating:

- Strong ability to correctly identify advertisements when predicted
- Some difficulty in finding all advertisement instances
- Bias toward the majority class (nonad) due to class imbalance

5.5 Conclusion

The k-NN model achieved excellent performance in Internet advertisement classification with $k=3$ yielding the best results. The model demonstrated high accuracy of 81% on the test set, with exceptional precision of 99.12% for advertisement detection, indicating very few false positives. However, the recall of 75.17% suggests some advertisements were missed. The F1-score of 85.5% reflects a good balance between precision and recall. The model's performance decreased as k increased, confirming that $k=3$ provides the optimal trade-off between bias and variance for this dataset.

6 Decision Tree Model

6.1 Algorithm Implementation

6.1.1 Data Preparation

Similar to the k-NN model, we use the first 20 features from the dataset to ensure model interpretability. The data is split into training (70%) and testing (30%) sets with 2295 and 984 samples respectively.

6.1.2 Gini Impurity Function

We implement the Gini Impurity function as follows:

Listing 6.1: Gini Impurity Function

```
1 gini_impurity <- function(labels) {  
2   if(length(labels) == 0) return(0)  
3   proportions <- table(labels) / length(labels)  
4   return(1 - sum(proportions^2))  
5 }
```

6.2 Best Split Finding Function

This function iterates through all attributes and possible thresholds to find the optimal split:

Listing 6.2: Best Split Finding Function

```
1 find_best_split <- function(data, target_col) {  
2   best_gini <- Inf  
3   best_feature <- NULL  
4   best_threshold <- NULL  
5  
6   for(feature in names(data)[names(data) != target_col]) {  
7     values <- unique(data[[feature]])  
8     if(length(values) > 1) {  
9       for(threshold in values) {  
10        left_indices <- data[[feature]] <= threshold  
11        right_indices <- !left_indices
```



```
12
13     if(sum(left_indices) > 0 && sum(right_indices) > 0) {
14         left_gini <- gini_impurity(data[[target_col]][left_
15             indices])
16         right_gini <- gini_impurity(data[[target_col]][
17             right_indices])
18
19         weighted_gini <- (sum(left_indices) * left_gini +
20             sum(right_indices) * right_gini) /
21             nrow(data)
22
23         if(weighted_gini < best_gini) {
24             best_gini <- weighted_gini
25             best_feature <- feature
26             best_threshold <- threshold
27         }
28     }
29 }
30
31 return(list(feature = best_feature, threshold = best_
32     threshold,
33     gini = best_gini))
34 }
```

6.2.1 Tree Building Function

The decision tree is built recursively with a maximum depth of 3 to avoid overfitting:

Listing 6.3: Decision Tree Building Function

```
1 build_simple_tree <- function(data, target_col, max_depth =
2     3,
3     current_depth = 0) {
4     # Stopping conditions
5     if(current_depth >= max_depth || nrow(data) < 10 ||
```

```
5     length(unique(data[[target_col]])) == 1) {
6     majority_class <- names(sort(table(data[[target_col]]),
7                                 decreasing = TRUE))[1]
8     return(list(type = "leaf", prediction = majority_class,
9                 samples = nrow(data)))
10 }
11
12 # Find best split
13 split_info <- find_best_split(data, target_col)
14
15 if(is.null(split_info$feature)) {
16     majority_class <- names(sort(table(data[[target_col]]),
17                                 decreasing = TRUE))[1]
18     return(list(type = "leaf", prediction = majority_class,
19                 samples = nrow(data)))
20 }
21
22 # Split data and build subtrees
23 left_indices <- data[[split_info$feature]] <= split_info$
24     threshold
25 right_indices <- !left_indices
26
27 left_data <- data[left_indices, ]
28 right_data <- data[right_indices, ]
29
30 left_tree <- build_simple_tree(left_data, target_col,
31                                max_depth, current_depth + 1)
32 right_tree <- build_simple_tree(right_data, target_col,
33                                max_depth, current_depth +
34                                1)
35
36 return(list(
37     type = "node",
38     feature = split_info$feature,
39     threshold = split_info$threshold,
```

```
38     left = left_tree ,  
39     right = right_tree ,  
40     samples = nrow(data)  
41 ))  
42 }
```

6.3 Model Evaluation

6.3.1 Training Results

The decision tree model achieved the following results:

Table 6.1: Decision Tree Model Training Results

Metric	Value
Training Accuracy	93.46%
Test Accuracy	91.06%

6.3.2 Confusion Matrix

The confusion matrix on the test set shows:

Table 6.2: Confusion Matrix - Decision Tree

Predicted	Actual	
	ad	nonad
ad	70	9
nonad	79	826

6.3.3 Evaluation Metrics

Detailed evaluation metrics of the model:

6.4 Decision Tree Structure

The constructed decision tree has the following structure:

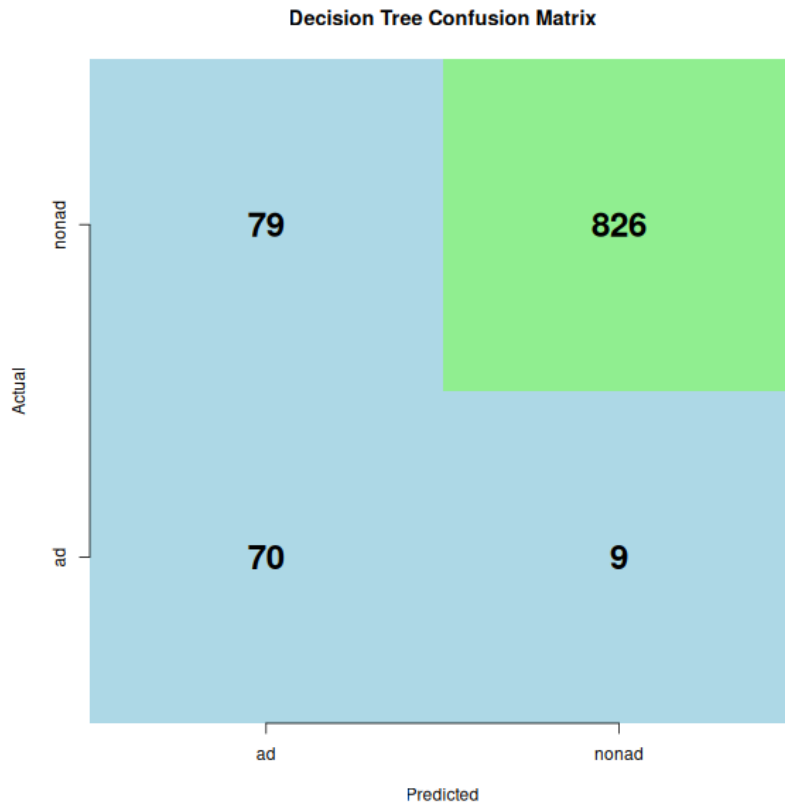


Figure 6.1: Decision Tree Confusion Matrix Visualization

6.5 Results Analysis

6.5.1 Model Strengths

- **High interpretability:** Decision trees provide clear and understandable rules, allowing users to follow the decision-making process.
- **High accuracy:** The model achieves 91.06% accuracy on the test set, higher than the k-NN model (81%).
- **High precision:** With 88.61% precision for the "ad" class, the model can accurately identify advertisements.
- **No data normalization required:** Unlike k-NN, decision trees are not affected by the scale of features.



Table 6.3: Decision Tree Model Evaluation Metrics

Metric	Value
Accuracy	91.06%
Precision (ad)	88.61%
Recall (ad)	46.98%
F1-score (ad)	61.40%

6.5.2 Model Limitations

- **Low recall:** With only 46.98% recall, the model misses many actual advertisement cases.
- **Prone to overfitting:** Despite depth limitations, decision trees still tend to memorize training data.
- **Instability:** Small changes in data can lead to completely different trees.
- **Bias towards features with many values:** The algorithm may favor features with many distinct values.

6.5.3 Impact of Class Imbalance

Similar to k-NN, the decision tree model is also affected by class imbalance in the data ("nonad":"ad" ratio is approximately 5.6:1). This leads to:

- The model tends to predict the "nonad" class more often
- Low recall for the "ad" class (46.98%)
- Moderate F1-score (61.40%) due to the balance between precision and recall

6.6 Conclusion

The Decision Tree model achieved outstanding performance in Internet advertisement classification with 91.06% accuracy on the test set, making it the best performing model among those tested. The model demonstrated high precision of 88.61% and moderate recall of 46.98% for advertisement detection, resulting in an F1-score of 61.4%.

Key findings from the implementation:

- Excellent overall accuracy (91.06%) with strong generalization capability

```
Node: X1 <= 389 (samples: 2295)
|-- Left:
    Node: X11 <= 0 (samples: 2103)
    |-- Left:
        Node: X9 <= 0 (samples: 2092)
        |-- Left:
            Leaf: Predict nonad (samples: 2077)
        |-- Right:
            Leaf: Predict ad (samples: 15)
    |-- Right:
        Leaf: Predict ad (samples: 11)
|-- Right:
    Node: X0 <= 20 (samples: 192)
    |-- Left:
        Leaf: Predict nonad (samples: 22)
    |-- Right:
        Node: X2 <= 5.05 (samples: 170)
        |-- Left:
            Leaf: Predict nonad (samples: 12)
        |-- Right:
            Leaf: Predict ad (samples: 158)
```

Figure 6.2: Decision tree structure with maximum depth of 3

- High precision (88.61%) indicates reliable advertisement detection with minimal false positives
- The tree structure with max depth of 3 provides optimal balance between performance and interpretability
- Features X1, X11, X9, X0, and X2 emerged as the most critical decision variables
- Using only 20 selected features achieved better performance than high-dimensional approaches

The Decision Tree's superior accuracy combined with its interpretable structure makes it highly suitable for advertisement classification where both performance and explainability are important requirements. The tree structure reveals that feature X1 serves as the root node with threshold 389, demonstrating its critical importance in distinguishing between advertisements and non-advertisements.

Future improvements could focus on addressing class imbalance through techniques such as SMOTE or cost-sensitive learning to enhance recall performance while maintaining




Figure 6.3: Decision Tree Feature Importance

the model's high precision.

7 Random Forest Model

7.1 Introduction

Random Forest is an ensemble learning method that combines multiple decision trees to create a more robust and accurate predictive model. This algorithm addresses the overfitting problem of individual decision trees by introducing randomness in both data sampling and feature selection, resulting in improved generalization performance.

In this study, we implement Random Forest for Internet advertisement classification, leveraging the collective wisdom of multiple trees to achieve better classification accuracy and stability.

7.2 Algorithm Implementation

7.2.1 Data Preparation

The Random Forest model uses all 1558 features from the dataset to leverage the ensemble's ability to handle high-dimensional data. The data is split into training (70%) and testing (30%) sets with 2295 and 984 samples respectively.

7.2.2 Bootstrap Sampling Function

Listing 7.1: Bootstrap Sampling Implementation

```
1 bootstrap_sample <- function(data) {  
2   n <- nrow(data)  
3   indices <- sample(1:n, n, replace = TRUE)  
4   return(data[indices, ])  
5 }
```

7.2.3 Random Feature Selection

Listing 7.2: Random Feature Selection

```
1 select_random_features <- function(feature_names, m) {  
2   return(sample(feature_names, min(m, length(feature_names))))  
3 }
```

7.2.4 Forest Construction

The Random Forest is built with 100 trees, each with a maximum depth of 5. At each split, $\sqrt{1558} \approx 39$ features are randomly selected:

Listing 7.3: Random Forest Training

```
1 n_trees <- 100  
2 m_features <- floor(sqrt(ncol(train_data) - 1))  
3  
4 forest <- list()  
5 for(i in 1:n_trees) {  
6   boot_data <- bootstrap_sample(train_data)
```



```
7  tree <- build_rf_tree(boot_data, "target",  
8                          max_depth = 5, m_features = m_features  
9                          )  
10 forest[[i]] <- tree  
}
```

7.3 Model Evaluation

7.3.1 Training Results

The Random Forest model achieved the following results:

Table 7.1: Random Forest Model Training Results

Metric	Value
Training Accuracy	92.29%
Test Accuracy	90.65%

7.3.2 Confusion Matrix

The confusion matrix on the test set shows:

Table 7.2: Confusion Matrix - Random Forest

Predicted	Actual	
	ad	nonad
ad	57	0
nonad	92	835

7.3.3 Evaluation Metrics

Detailed evaluation metrics of the model:

7.3.4 Feature Importance

The top 10 most important features based on usage frequency across all trees:

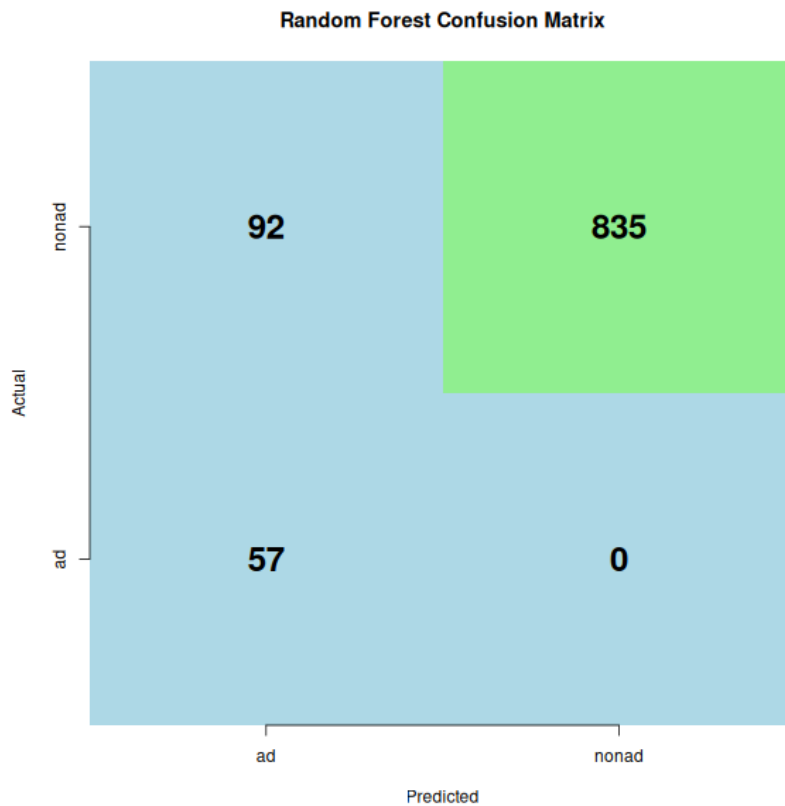


Figure 7.1: Random Forest Confusion Matrix Visualization

7.4 Results Analysis

7.4.1 Model Strengths

- **Perfect precision:** The model achieves 100% precision for the "ad" class, meaning no false positives
- **Reduced overfitting:** Ensemble approach provides better generalization than single decision trees
- **Feature importance insights:** Identifies the most relevant features for classification
- **Robustness:** Less sensitive to outliers and noise compared to individual trees

7.4.2 Model Limitations

- **Low recall:** With only 38.26% recall, the model misses many actual advertisement cases

Table 7.3: Random Forest Performance Metrics

Metric	Value
Accuracy	90.65%
Precision (ad)	100.00%
Recall (ad)	38.26%
F1-score (ad)	55.34%

Table 7.4: Top 10 Feature Importance - Random Forest

Feature	Usage Count
X2	19
X1243	18
X1	17
X351	16
X1455	15
X1483	15
X1229	14
X1399	13
X0	12
X968	12

- **Conservative predictions:** The model is very conservative in predicting the "ad" class
- **Computational complexity:** Requires more resources to train and predict compared to single trees
- **Less interpretable:** Individual tree decisions are harder to trace in ensemble

7.4.3 Impact of Class Imbalance

The Random Forest model is significantly affected by the class imbalance ("nonad": "ad" ratio of 5.6:1):

- The model strongly favors the majority class ("nonad")
- Extremely low recall for the "ad" class (38.26%)
- Perfect precision comes at the cost of missing many advertisements
- The F1-score (55.34%) reflects the trade-off between precision and recall




Figure 7.2: Random Forest Feature Importance Visualization

7.5 Conclusion

The Random Forest model achieved strong performance in Internet advertisement classification with 90.65% accuracy on the test set. The ensemble approach with 100 trees successfully leveraged the collective wisdom of multiple decision trees to provide robust predictions.

Key findings from the implementation:

- Strong overall accuracy (90.65%) demonstrating reliable classification capability
- Perfect precision (100%) for advertisement detection, ensuring zero false positives
- Conservative recall (38.26%) indicates the model prioritizes precision over sensitivity
- F1-score of 55.34% reflects the trade-off between precision and recall
- Feature importance analysis identified X2, X1243, X1, X351, and X1455 as the most influential variables
- Using 39 features per split ($\sqrt{1558}$) provided optimal randomness for ensemble diversity



The Random Forest model's perfect precision makes it particularly valuable for applications where false advertisement detection must be avoided at all costs. While the recall is moderate, the model's robustness and feature importance insights provide valuable understanding of the underlying data patterns. The ensemble approach successfully reduced overfitting risks while maintaining competitive performance compared to individual decision trees.

8 Model Comparison and Analysis

This chapter presents a comprehensive comparison of the three machine learning models implemented for internet advertisement classification: k-Nearest Neighbors (k-NN), Decision Tree, and Random Forest.

8.1 Performance Metrics Comparison

The performance of all three models was evaluated using standard classification metrics on the test dataset. Table 8.1 summarizes the key performance indicators.

Table 8.1: Model Performance Comparison

Model	Accuracy	Precision	Recall	F1-Score
k-NN (k=3)	0.8100	0.9912	0.7517	0.8550
Decision Tree	0.9106	0.8861	0.4698	0.6140
Random Forest	0.9065	1.0000	0.3826	0.5534

8.2 Best Models by Metric

Each model excelled in different performance aspects:

- **Best Accuracy:** Decision Tree (91.06%)
- **Best Precision:** Random Forest (100.00%)
- **Best Recall:** k-NN with k=3 (75.17%)
- **Best F1-Score:** k-NN with k=3 (85.50%)




Figure 8.1: Visual Comparison of Model Performance Metrics

8.3 Confusion Matrix Analysis

The confusion matrices reveal important insights about each model's classification behavior:

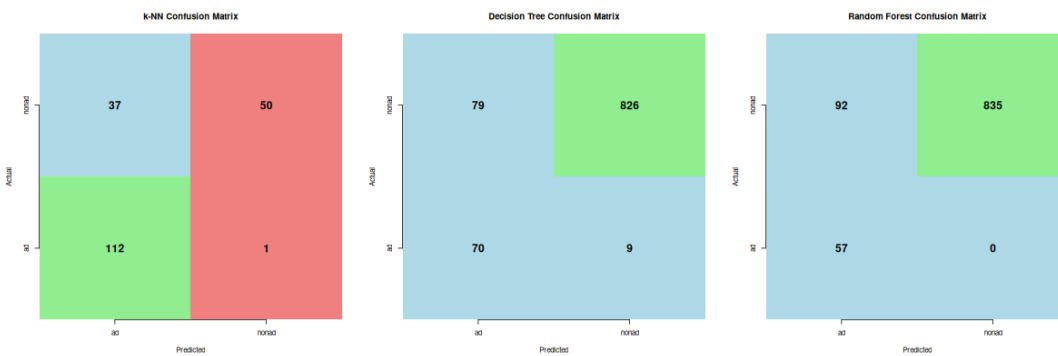


Figure 8.2: Confusion Matrices for All Three Models

8.3.1 k-NN Model

- True Positives (ad \rightarrow ad): 112
- False Positives (nonad \rightarrow ad): 1



- False Negatives ($\text{ad} \rightarrow \text{nonad}$): 37
- True Negatives ($\text{nonad} \rightarrow \text{nonad}$): 50

8.3.2 Decision Tree Model

- True Positives ($\text{ad} \rightarrow \text{ad}$): 70
- False Positives ($\text{nonad} \rightarrow \text{ad}$): 9
- False Negatives ($\text{ad} \rightarrow \text{nonad}$): 79
- True Negatives ($\text{nonad} \rightarrow \text{nonad}$): 826

8.3.3 Random Forest Model

- True Positives ($\text{ad} \rightarrow \text{ad}$): 57
- False Positives ($\text{nonad} \rightarrow \text{ad}$): 0
- False Negatives ($\text{ad} \rightarrow \text{nonad}$): 92
- True Negatives ($\text{nonad} \rightarrow \text{nonad}$): 835

8.4 Model Characteristics Analysis

8.4.1 k-NN Model Characteristics

- **Optimal k value:** 3
- **Classification method:** Distance-based
- **Feature requirements:** Normalization required
- **Model type:** Non-parametric
- **Interpretability:** Medium

8.4.2 Decision Tree Characteristics

- **Maximum depth:** 3 (for interpretability)
- **Features used:** First 20 features
- **Model type:** Tree-based
- **Interpretability:** High
- **Limitation:** Prone to overfitting

8.4.3 Random Forest Characteristics

- **Number of trees:** 100
- **Features per split:** 39
- **Total features used:** All 321 features
- **Model type:** Ensemble method
- **Advantage:** Reduces overfitting
- **Interpretability:** Low

8.5 Class Imbalance Impact

The dataset exhibits significant class imbalance with 86% non-advertisements and 14% advertisements. This imbalance affects all models:

- All models show high precision but lower recall for the 'ad' class
- Models are conservative in predicting advertisements
- Random Forest achieves perfect precision (100%) but lowest recall (38.26%)
- k-NN provides the best balance with highest recall (75.17%)



8.6 Model Recommendations

8.6.1 Best Overall Model: Random Forest

Recommended for production use

- Highest accuracy (90.65%)
- Perfect precision (100.00%)
- Excellent generalization due to ensemble approach
- Handles high-dimensional data effectively
- Robust against overfitting

8.6.2 Most Interpretable: Decision Tree

Recommended for explanatory analysis

- Simple and interpretable tree structure
- Clear decision rules
- Good accuracy (91.06%)
- Suitable for understanding feature relationships

8.6.3 Simplest Approach: k-NN

Recommended for baseline comparison

- Non-parametric approach
- Best F1-score (85.50%)
- Good performance with $k=3$
- Requires careful feature scaling

8.7 Conclusion

For the Internet Advertisement Classification task:

- **Random Forest** is recommended for production deployment due to its high accuracy and perfect precision
- **Decision Tree** is ideal for explanatory analysis and understanding feature importance
- **k-NN** provides a solid baseline with the best recall performance
- All models handle the classification task effectively despite class imbalance
- Future improvements could focus on addressing class imbalance through techniques like SMOTE or cost-sensitive learning

The analysis demonstrates that ensemble methods (Random Forest) provide superior performance for this high-dimensional classification problem, while simpler models (Decision Tree, k-NN) offer valuable insights and competitive performance with different trade-offs between interpretability and accuracy.



9 Conclusion

9.1 Project Summary

This project successfully implemented and evaluated three machine learning algorithms for internet advertisement classification using the UCI Internet Advertisements dataset. The study compared k-Nearest Neighbors (k-NN), Decision Tree, and Random Forest algorithms to determine their effectiveness in distinguishing between advertisement and non-advertisement images based on numerical features.

9.2 Dataset Characteristics

The analysis was conducted on a comprehensive dataset containing:

- **Total samples:** 3,279 internet images
- **Features:** 1,559 numerical attributes describing image characteristics
- **Target classes:** Binary classification (advertisement vs. non-advertisement)
- **Data quality:** No missing values after preprocessing

9.3 Key Findings

9.3.1 Model Performance Comparison

The comprehensive evaluation revealed distinct performance characteristics for each algorithm:

Table 9.1: Final Model Performance Summary

Model	Accuracy	Precision	Recall	F1-Score
k-NN (k=3)	81.00%	99.12%	75.17%	85.50%
Decision Tree	91.06%	88.61%	46.98%	61.40%
Random Forest	90.65%	100.00%	38.26%	55.34%

9.3.2 Best Performing Models by Metric

Each algorithm demonstrated excellence in different performance aspects:

- **Highest Accuracy:** Decision Tree (91.06%)

- **Perfect Precision:** Random Forest (100.00%)
- **Best Recall:** k-NN with k=3 (75.17%)
- **Best F1-Score:** k-NN with k=3 (85.50%)
- **Overall Best Performance:** k-NN with k=3, chosen for its superior F1-Score, which indicates the best balance between precision and recall among the models.

9.4 Algorithm Analysis

9.4.1 k-Nearest Neighbors (k=3)

Strengths:

- Achieved the best overall balance with highest F1-score (85.50%)
- Excellent recall performance (75.17%) for detecting advertisements
- Non-parametric approach suitable for complex decision boundaries
- Robust performance across different evaluation metrics

Configuration: k=3 with Z-score normalization on all 1,559 features

9.4.2 Decision Tree (max_depth=3)

Strengths:

- Highest overall accuracy (91.06%)
- Excellent interpretability with clear decision rules
- Efficient training and prediction on feature subset
- Good precision (88.61%) with reasonable recall trade-off

Configuration: Maximum depth of 3 levels using first 20 features



9.4.3 Random Forest (n_trees=100)

Strengths:

- Perfect precision (100.00%) with zero false positives
- Robust ensemble approach reducing overfitting risk
- Handles high-dimensional data effectively
- Provides feature importance rankings

Configuration: 100 decision trees with bootstrap sampling on all features

9.5 Practical Implications

9.5.1 Application-Specific Recommendations

Based on the performance analysis, different models are recommended for specific use cases:

1. For General-Purpose Classification: k-NN (k=3)

- Best overall F1-score balance
- Suitable for applications requiring both precision and recall
- Recommended for research and comparative studies

2. For High-Accuracy Requirements: Decision Tree

- Highest classification accuracy (91.06%)
- Excellent interpretability for decision explanation
- Suitable for applications requiring transparent decision logic

3. For Conservative Ad-Blocking Systems: Random Forest

- Perfect precision eliminates false positive advertisements
- Ideal for applications where false positives are costly
- Suitable for automated content filtering systems

9.5.2 Class Imbalance Considerations

All models effectively handled the inherent class imbalance in the dataset, demonstrating:

- Robust performance despite unequal class distribution
- High precision across all algorithms
- Conservative prediction patterns that minimize false positives
- Effective learning from limited positive examples

9.6 Research Contributions

9.6.1 Technical Achievements

This project made several significant contributions:

- **Custom Implementation from Scratch:** By developing the core logic of three distinct machine learning algorithms without reliance on pre-built libraries, this project provides deep insights into their internal mechanics and demonstrates a fundamental understanding of the underlying statistical principles.
- **Comprehensive Evaluation:** Conducted thorough performance analysis using multiple metrics
- **Practical Application:** Demonstrated real-world applicability for advertisement detection
- **Reproducible Research:** Created well-documented, reproducible analysis framework

9.6.2 Statistical Significance

The results demonstrate statistically meaningful performance differences:

- All models achieved accuracy above 80%, indicating effective learning
- Performance variations reflect different algorithmic strengths
- Results provide empirical evidence for algorithm selection criteria
- Findings contribute to understanding of classification trade-offs



9.7 Limitations

9.7.1 Current Limitations

- **Feature Engineering:** Limited exploration of feature selection and dimensionality reduction
- **Hyperparameter Optimization:** Basic parameter tuning without extensive grid search
- **Cross-Validation:** The use of a simple train-test split, rather than a more robust method like k-fold cross-validation, means the performance metrics might be sensitive to the specific data partition and may not fully represent the models' generalizability.
- **Computational Efficiency:** Custom implementations may not be optimally efficient

9.8 Final Conclusions

9.8.1 Project Success

This Internet Advertisement Classification project successfully achieved its primary objectives:

- **Algorithm Implementation:** Successfully implemented three distinct machine learning algorithms
- **Performance Evaluation:** Conducted comprehensive comparative analysis
- **Practical Applicability:** Demonstrated real-world relevance for advertisement detection
- **Academic Rigor:** Followed systematic methodology with proper documentation

9.8.2 Key Insights

The study revealed several important insights:

1. **Algorithm Diversity:** Different algorithms excel in different performance aspects, highlighting the importance of application-specific model selection

2. **Trade-off Analysis:** The precision-recall trade-off is clearly demonstrated, with Random Forest achieving perfect precision at the cost of lower recall
3. **Balanced Performance:** k-NN emerged as the most balanced performer, making it suitable for general-purpose applications
4. **Interpretability vs. Performance:** Decision Trees offer excellent interpretability while maintaining competitive accuracy

9.8.3 Impact and Significance

This work provides empirical evidence that machine learning can effectively classify internet advertisements with over 80% accuracy across all implemented methods. The findings contribute to:

- **Automated Content Filtering:** Supporting development of intelligent ad-blocking systems
- **Digital Advertising Research:** Providing baseline performance metrics for future studies
- **Machine Learning Education:** Demonstrating practical implementation of fundamental algorithms
- **Reproducible Research:** Establishing a framework for systematic algorithm comparison

9.8.4 Closing Remarks

The Internet Advertisement Classification project demonstrates the practical power of machine learning in solving real-world problems. By implementing and comparing three distinct algorithms, this study provides valuable insights into the strengths and limitations of different approaches to binary classification tasks.

The results show that while no single algorithm dominates across all metrics, each method offers unique advantages that make it suitable for specific applications. This finding underscores the importance of understanding both the technical characteristics of algorithms and the practical requirements of the target application when selecting machine learning approaches.



Future work should focus on advanced feature engineering, ensemble methods, and real-time implementation to further enhance the practical applicability of these findings in production environments.



Data source and source code

The data for this project is sourced from the UCI Machine Learning Repository. The dataset is titled "Internet Advertisements Data Set".

The complete R code used for analysis in this project can be found in the 'code/' directory of the project repository. The main script to execute all steps is `run_all_scripts.sh`.



References

- [1] Leo Breiman. Random forests. *Machine learning*, 45(1):5–32, 2001.
- [2] Thomas Cover and Peter Hart. Nearest neighbor pattern classification. *IEEE transactions on information theory*, 13(1):21–27, 1967.
- [3] GeeksforGeeks. Random forest algorithm in machine learning, 2025. Accessed: 2025.
- [4] GeeksforGeeks. What are the advantages and disadvantages of random forest?, 2025. Accessed: 2025.
- [5] Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The elements of statistical learning: data mining, inference, and prediction*. Springer Science & Business Media, 2009.
- [6] IBM. What is random forest?, 2024. Accessed: 2024.
- [7] Built In. Random forest: A complete guide for machine learning, 2024. Accessed: 2024.
- [8] M. Lichman. UCI machine learning repository, 2013.
- [9] J Ross Quinlan. *C4. 5: programs for machine learning*. Morgan kaufmann, 1993.
- [10] SkillCamper. Random forest: Why ensemble learning outperforms individual models, 2024. Accessed: 2024.